

A Study of LSM-Tree

林金河

Applications of LSM-Tree

❑ NoSQL

- ❑ Bigtable
- ❑ HBase
- ❑ Cassandra
- ❑ Scylla
- ❑ MongoDB



❑ Storage Engine

- ❑ LevelDB
- ❑ RocksDB
- ❑ MyRocks



❑ NewSQL


- ❑ TiDB(TiKV)
- ❑ CockroachDB



Beginning of LSM-Tree

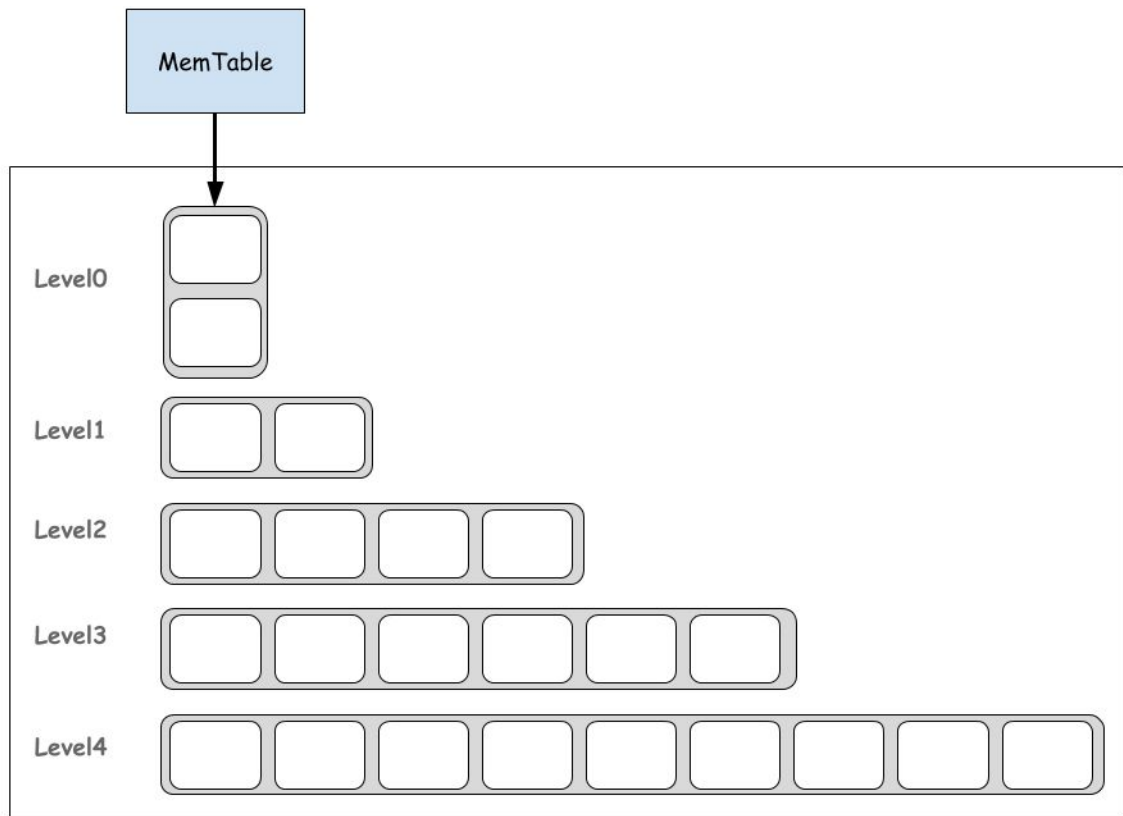
- ❑ 1996: The Log-Structured Merge-Tree (LSM-Tree)
 - ❑ Memory is small
 - ❑ Disk is slow for random access(r/w)
 - ❑ Originally designed for fast-growing History table
 - ❑ Data and indexes
 - ❑ Write heavy
 - ❑ Read sparse

Beginning of LSM-Tree



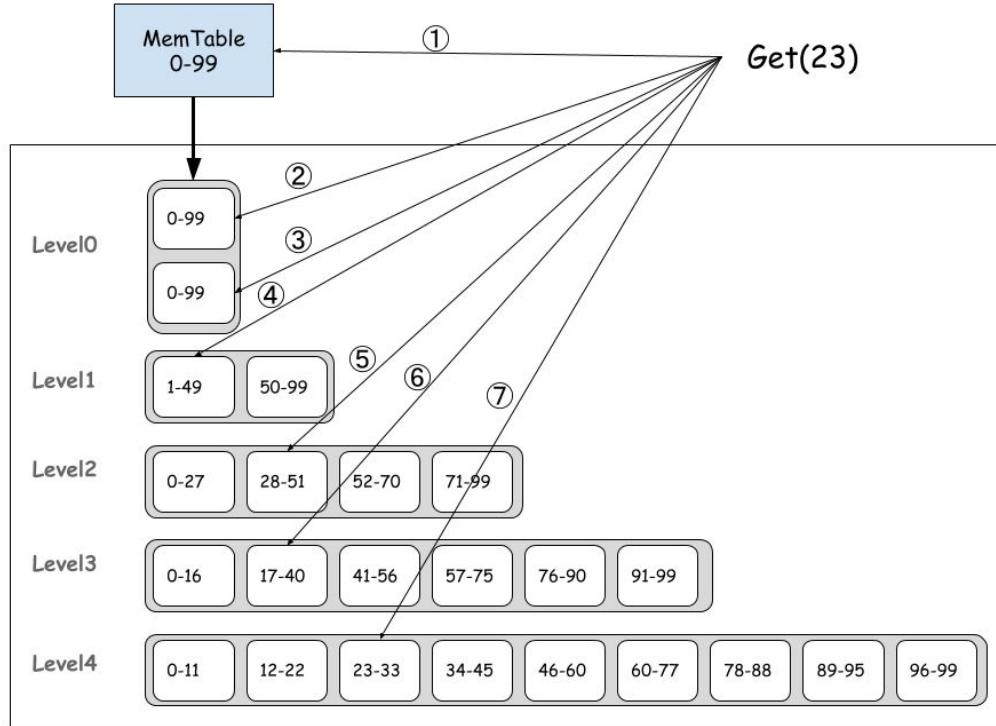
- ❑ Out-of-place update
- ❑ Optimized for write
- ❑ Sacrifice read
- ❑ Not optimized for space
- ❑ Require data reorganization (merge/compaction)

Modern Structure of LSM-Tree




- ❑ Read
 - ❑ Point Query
 - ❑ Range Query
- ❑ Write(Insert/Delete/Update)
- ❑ Compaction
 - ❑ Leveled
 - ❑ Tired

Point Query



- ❑ Returns immediately when something found
- ❑ Read amplification, worst-case
I/O: $2 * (N - 1 + \text{files num of level-0})$
- ❑ Optimization
 - ❑ Page cache/Block cache
 - ❑ Bloom filter

Bloom Filter

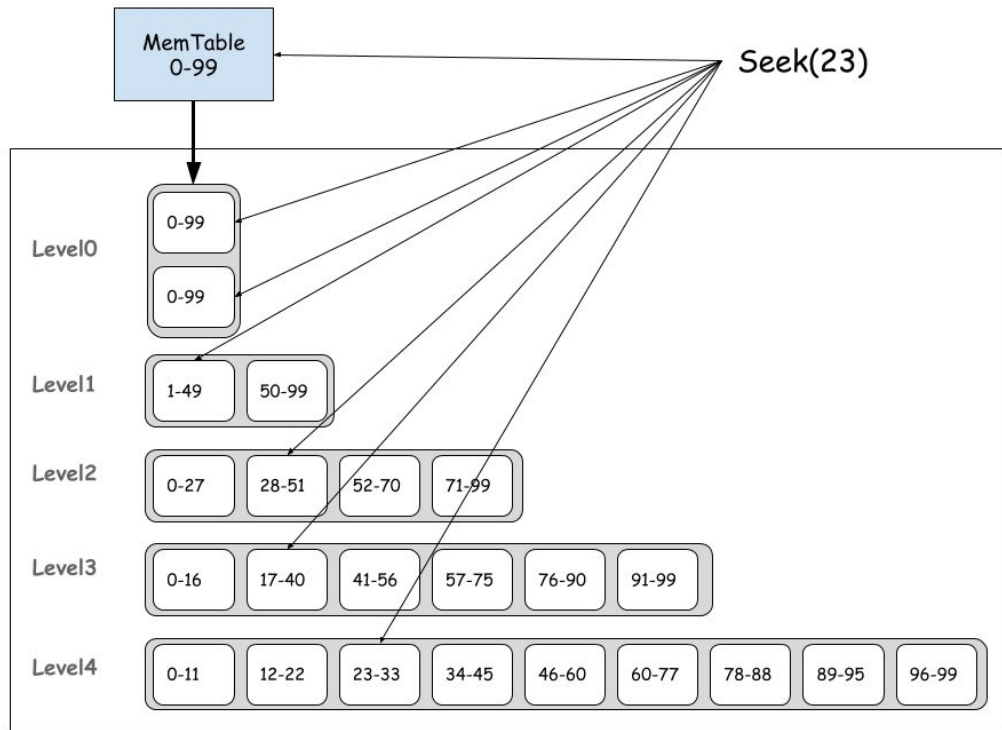


$$\frac{m}{n} = -\frac{\log_2 p}{\ln 2} \approx -1.44 \log_2 p$$

$$k = -\frac{\ln p}{\ln 2} = -\log_2 p$$

- ❑ Simple
- ❑ Space-efficient
- ❑ False positive is possible, but false negative is impossible
- ❑ Range query is **not** supported

Range Query



```
SELECT * FROM t WHERE key >= 23 AND  
key < 40;
```

```
for (itr->Seek(23); itr->Valid(); itr->Next()) {  
    if (itr->key() < 40) {  
        ...  
    } else ...  
}
```

- ❑ **Must seek every sorted run**
- ❑ Bloom filter **not** support range query
- ❑ Optimization
 - ❑ Parallel Seeks
 - ❑ Prefix bloom filter(RocksDB)
 - ❑ SuRF (SIGMOD 2018)

- ❑ SuRF: Practical Range Query Filtering with Fast Succinct Tries
- ❑ Fast Succinct Tries

-
- LOUDS:** 110 10 110 1110 110 110 0 10 0 0 0 10 0 0 0
 0 1 2 3 4 5 6 7 8 9 A B C D E

LOUDS-Dense

- D-Labels: f st a o r
- D-HasChild: [Indicators for child presence]
- D-ISPrefixKey: [Indicators for prefix keys]
- D-Values: v₁, v₂, 0

LOUDS-Sparse


- S-Labels: r s t p y i y \$ t e p
- S-HasChild: 0 1 0 0 0 1 0 0 0 0 0
- S-LOUDS: 1 0 0 1 0 1 0 1 0 1 0
- S-Values: v₃ v₄ v₅ v₆ v₇ v₈ v₉ v₁₀ v₁₁

Keys stored: f, fa, fas, fast, fat, s, top, toy, trie, trip, try

Figure 2: LOUDS-DS Encoded Trie – The \$ symbol represents the character whose ASCII number is 0xFF. It is used to indicate the situation where a prefix string leading to a node is also a valid key.



SuRF: Succinct Range Filter

- ❑ Trie \Rightarrow SuRF-Base \Rightarrow SuRF-Hash/SuRF-Real
- ❑ Example: ("SIGAI", "SIGMOD", "SIGOPS")




Trie

SuRF-Base stores the minimum-length key prefixes such that it can uniquely identify each key. Specifically, SuRF-Base only stores an additional byte for each key beyond the shared prefixes.



SuRF-Hash adds a few hash bits per key to SuRF-Base to reduce its FPR. The extra bits in SuRF-Hash do not help range queries because they do not provide ordering information on keys.



SuRF-Real stores the n key bits immediately following the stored prefix of a key. Both point and range queries benefit from the real suffix bits to reduce false positives.

SuRF: Succinct Range Filter

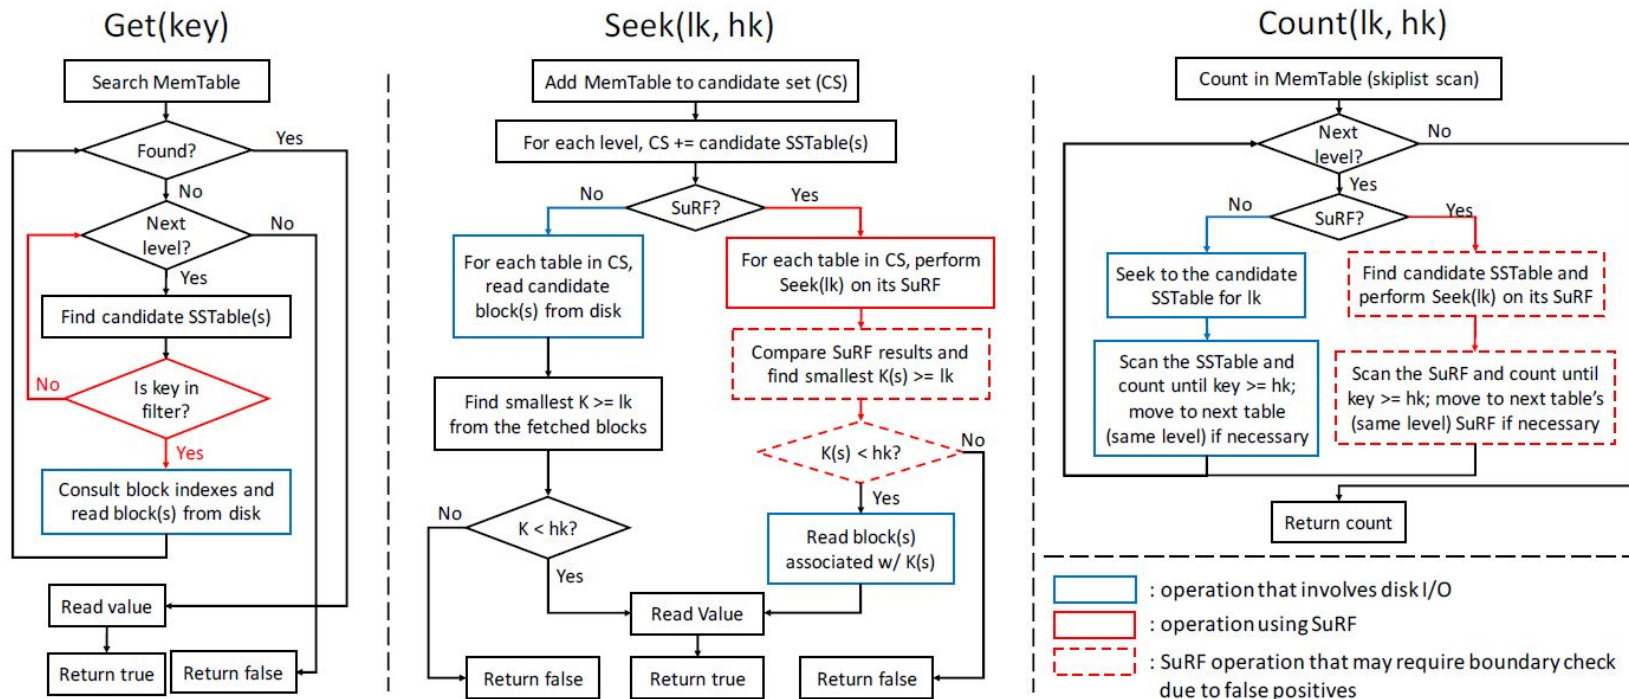

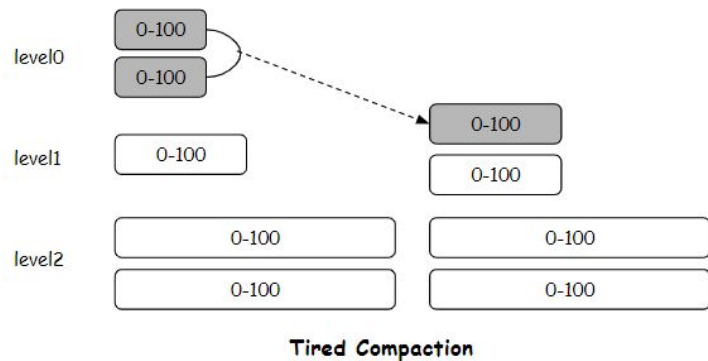


Figure 11: Execution paths for Get, Seek, and Count in RocksDB

Read Operation Summary

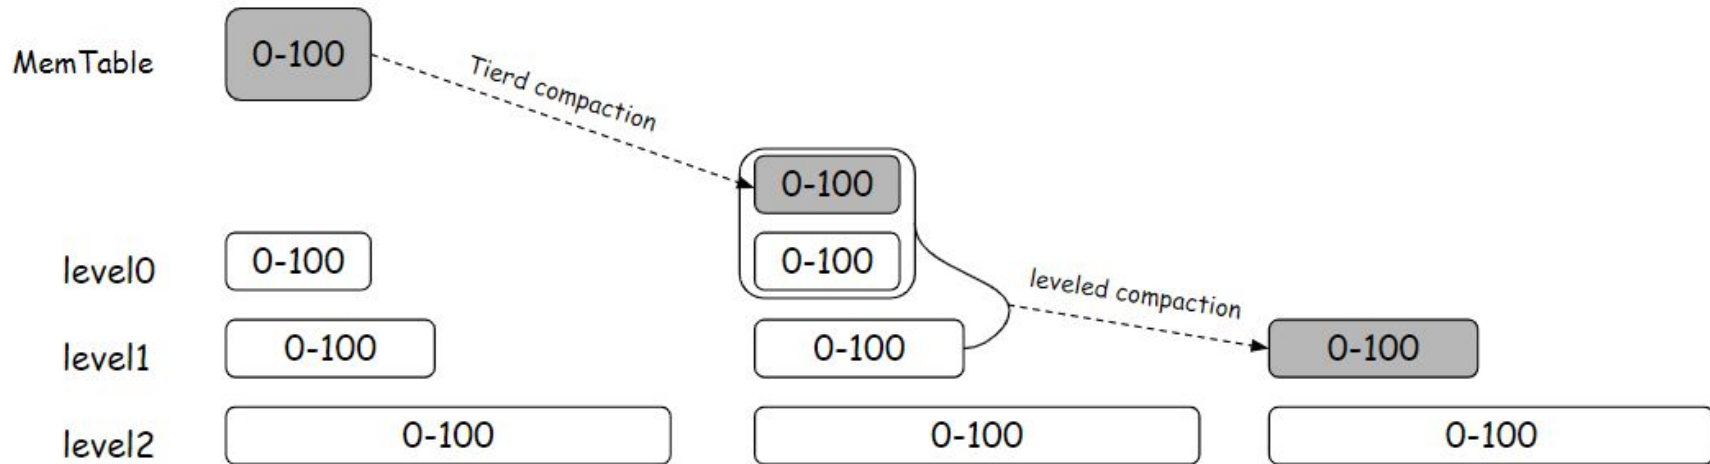
- ❑ Read amplification \Rightarrow Use filters to reduce unnecessary I/O
- ❑ Space-efficient
- ❑ Lookup performance
- ❑ Low false positive, no false negative
- ❑ Bloom filter/Prefix bloom filter
- ❑ SuRF
- ❑ Others: Cuckoo Filter...

Compaction - Tiered vs Leveled



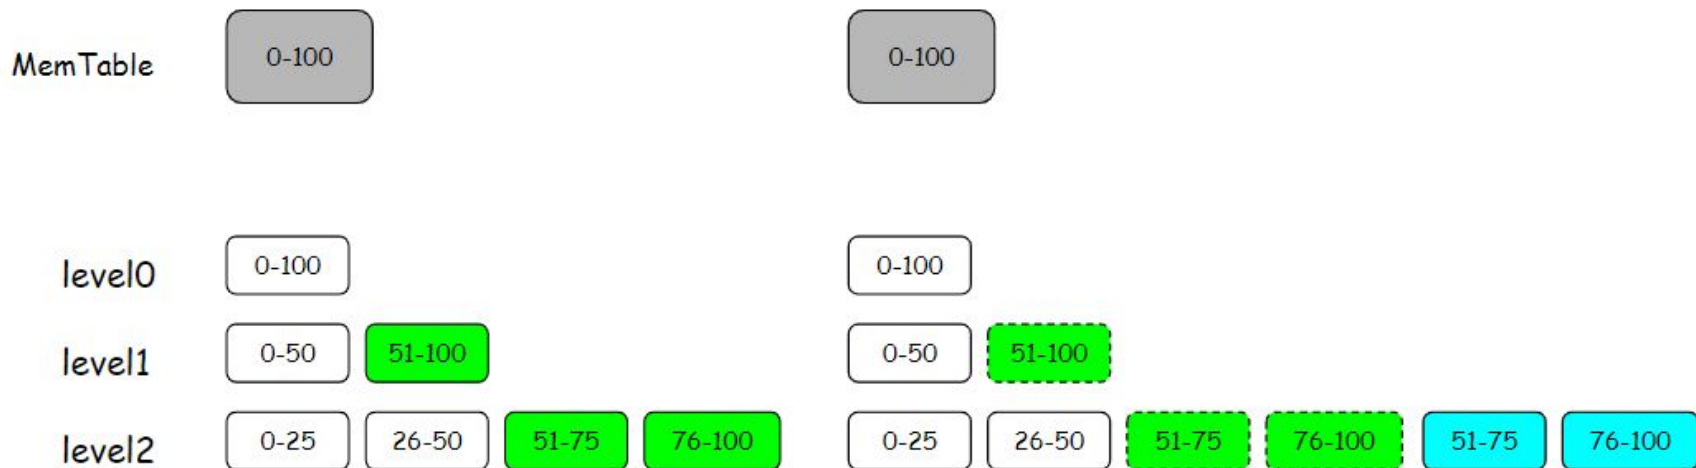
- ❑ Each level has **N sorted runs** (overlapped).
 - ❑ Compaction merges all sorted runs in one level to create a new sorted run in the next level.
 - ❑ Minimizes write amplification at the cost of read and space amplification.
-
- ❑ Each level is **one sorted run**.
 - ❑ Compaction into L_n merges data from L_{n-1} into L_n .
 - ❑ Compaction into L_n rewrites data that was previously merged into L_n .
 - ❑ Minimizes space amplification and read amplification at the cost of write amplification.

Tiered + Leveled



- ❑ Less write amplification than leveled and less space amplification than tiered.
- ❑ More read amplification than leveled and more write amplification than tiered.
- ❑ It is flexible about the level at which the LSM tree switches from tiered to leveled.

Tiered + Leveled + Partition




❑ Default compaction of LevelDB and RocksDB.

Problems of compaction

- ❑ Write amplification
 - ❑ System jitter(CPU, I/O)
 - ❑ SSD wear out
- ❑ Invalid block cache/page cache
- ❑ Compaction rate is a problem
 - ❑ Too fast - write amplification
 - ❑ Too slow - read amplification and space amplification.

Multi-Thread Compaction



Pipelined Compaction

❑ Pipelined Compaction for the LSM-tree

❑ The compaction procedure:

❑ **Step 1: Read data blocks.**

❑ Step2: Checksum.

❑ Step3: Decompress.

❑ Step4: Merge sort.

❑ Step5: Compress.

❑ Step6: Re-checksum.

❑ **Step7: Write to the disk.**

❑ Read → Compute → Write

❑ It is difficult to divide the stages evenly.

❑ If data blocks must flow through multiple processors, it will result in low CPU cache performance. Let S2~S6 as one stage will be more CPU cache friendly.


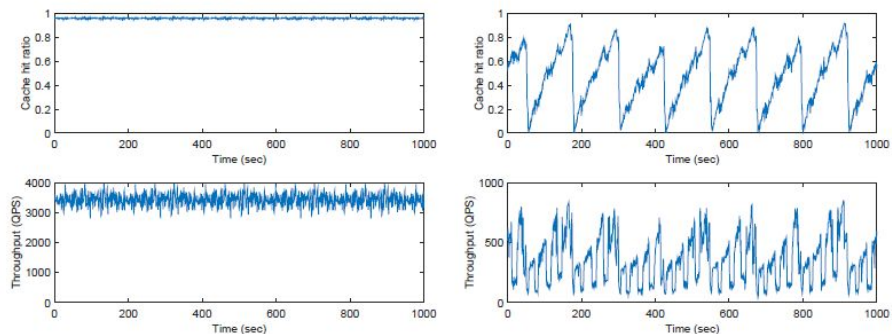


Fig. 4. Ideal Pipelined Compaction Procedure: *Step 1* and *Step 7* are scheduled on disk, and all the other steps are scheduled on CPU.

Compaction Buffer

- ❑ Re-enabling high-speed caching for LSM-trees
- ❑ LSbM-tree: Re-enabling Buffer Caching in Data Management for Mixed Reads and Writes



(a) Read only workload

(b) Read and write workload

Fig. 2. The failure of caching in LSM-tree

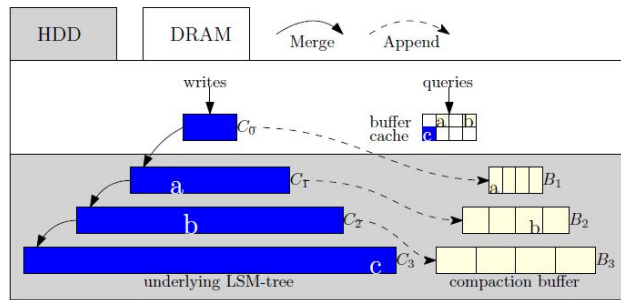


Fig. 5: The basic structure of a LSbM-tree





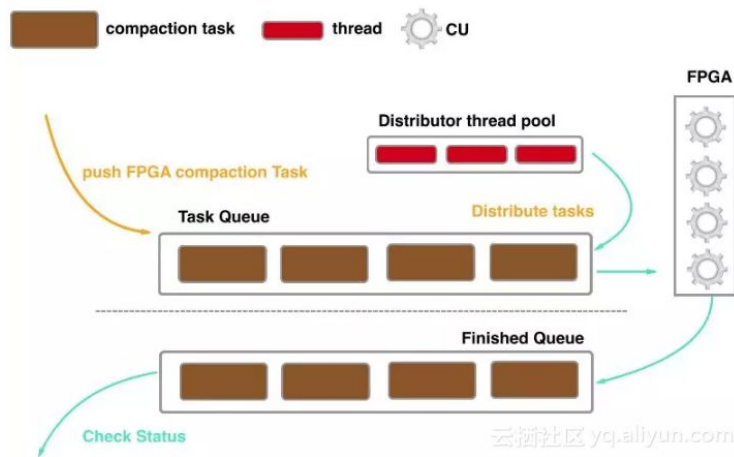
Fig. 6: An illustration of the buffered merge

Coprocessor

- ❑ Co-KV: A Collaborative Key-Value Store Using Near-Data Processing to Improve Compaction for the LSM-tree
- ❑ 当数据库遇见FPGA:X-DB异构计算如何实现百万级TPS？

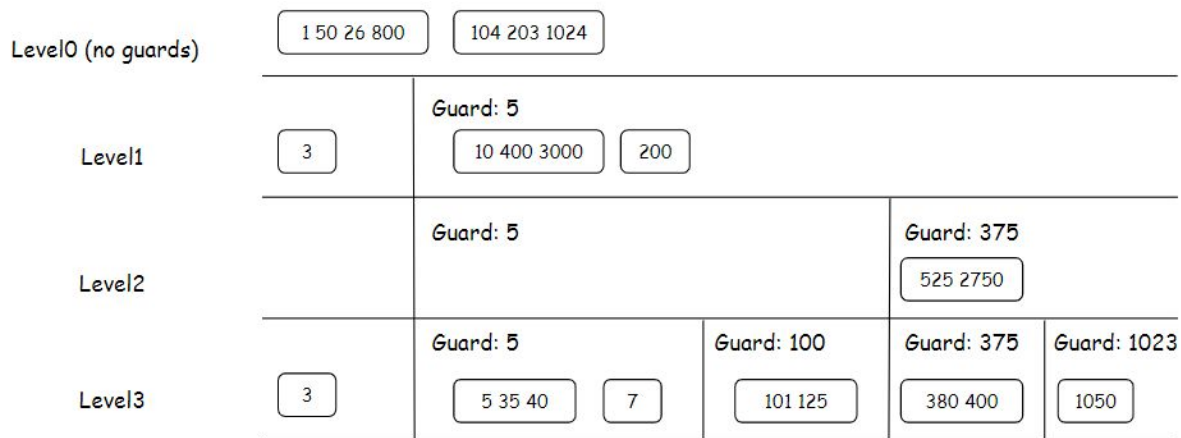


(b) Hybrid design, FPGA is integrated as a real co-processor.



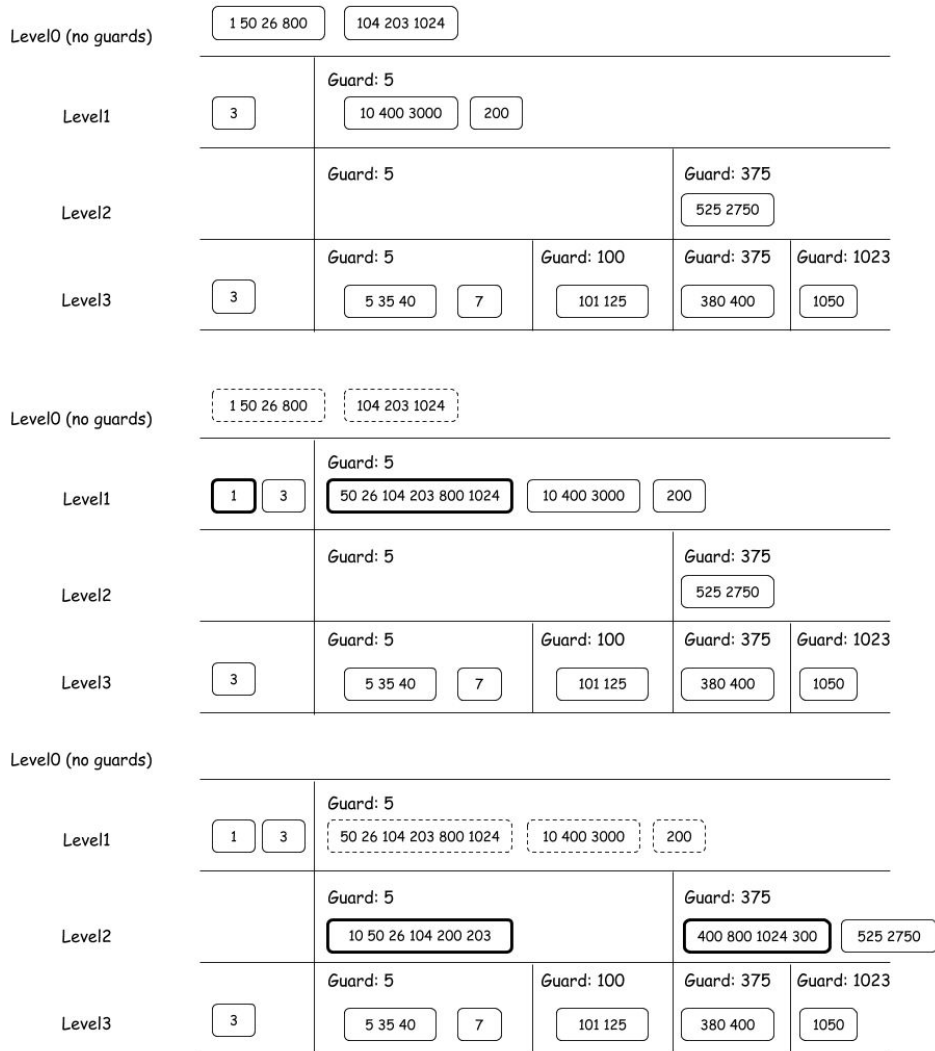
PebblesDB

- ❑ PebblesDB: Building Key-Value Stores using Fragmented Log-Structured Merge Trees
- ❑ **Guards** divide the key space (for that level) into disjoint units
- ❑ Less write amplification, more read amplification and space amplification.




PebblesDB Compaction

❑ Similar to Tired + Partition



WiscKey

- ❑ **WiscKey: Separating Keys from Values in SSD-conscious Storage**
- ❑ Compaction = sorting + garbage collection
- ❑ Only keys are required to be sorted
- ❑ Keys are usually smaller than values
- ❑ Key-Value separation \Rightarrow Decouple sorting and garbage collection



WiscKey

- ❑ Key-Value separation's challenges and optimizations
 - ❑ An additional I/O may be required for each query
 - ❑ WiscKey's LSM-Tree is small and can be easily cached in memory
 - ❑ Range queries require random I/O
 - ❑ Prefetch: leverages the parallel I/O characteristic of SSD
 - ❑ Garbage Collection
 - ❑ Invalid keys are reclaimed by compaction
 - ❑ Value log needs a special garbage collector
 - ❑ Crash Consistency
 - ❑ Atomicity of inserted key-value pairs is complicated
 - ❑ Combine value log and WAL

Summary

- ❑ Designing Access Methods: The RUM Conjecture
 - ❑ Read, Update, Memory - Optimize Two at the Expense of the Third
 - ❑ Read Amplification vs Write Amplification vs Space Amplification

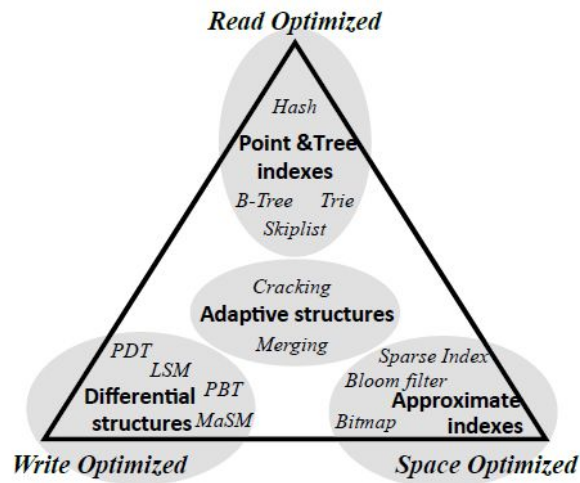


Figure 1: Popular data structures in the RUM space.

Q&A